

Import libraries

```
In [1]: import os
import cv2
import tensorflow as tf
import numpy as np
from typing import List
from matplotlib import pyplot as plt
import imageio
```

Build data loading functions

```
In [2]: def load_video(path:str) -> List[float]:

    cap = cv2.VideoCapture(path)
    frames = []
    for _ in range(int(cap.get(cv2.CAP_PROP_FRAME_COUNT))):
        ret, frame = cap.read()
        frame = tf.image.rgb_to_grayscale(frame)
        frames.append(frame[190:236,80:220,:])
    cap.release()

    mean = tf.math.reduce_mean(frames)
    std = tf.math.reduce_std(tf.cast(frames, tf.float32))
    return tf.cast((frames - mean), tf.float32) / std
```

```
In [3]: vocab = [x for x in "abcdefghijklmnopqrstuvwxyz?!123456789 "]

char_to_num = tf.keras.layers.StringLookup(vocabulary=vocab, oov_token="")
num_to_char = tf.keras.layers.StringLookup(
    vocabulary=char_to_num.get_vocabulary(), oov_token="", invert=True
)

print(
    f"The vocabulary is: {char_to_num.get_vocabulary()} "
    f"(size ={char_to_num.vocabulary_size()})"
)
```

The vocabulary is: ['', 'a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z', '!', '?', '1', '2', '3', '4', '5', '6', '7', '8', '9', ' '] (size =40)

```
In [4]: def load_alignments(path:str) -> List[str]:
    with open(path, 'r') as f:
        lines = f.readlines()
    tokens = []
    for line in lines:
        line = line.split()
        if line[2] != 'sil':
            tokens = [*tokens, ' ', line[2]]
    return char_to_num(tf.reshape(tf.strings.unicode_split(tokens, input_encoding='
```

```
In [5]: def load_data(path: str):
        path = bytes.decode(path.numpy())
        #file_name = path.split('/')[ -1].split('.')[0]
        # File name splitting for windows
        file_name = path.split('\\')[ -1].split('.')[0]
        video_path = os.path.join('data', 's1', f'{file_name}.mpg')
        alignment_path = os.path.join('data', 'alignments', 's1', f'{file_name}.align')
        frames = load_video(video_path)
        alignments = load_alignments(alignment_path)

        return frames, alignments
```

```
In [6]: def mappable_function(path:str) ->List[str]:
        result = tf.py_function(load_data, [path], (tf.float32, tf.int64))
        return result
```

Prepare testing dataset

```
In [7]: data = tf.data.Dataset.list_files('./data/s1/*.mpg')
        data = data.shuffle(500, reshuffle_each_iteration=False)
        data = data.map(mappable_function)
        data = data.padded_batch(2, padded_shapes=([75, None, None, None], [40]))
        data = data.prefetch(tf.data.AUTOTUNE)
        test = data.skip(450)
```

Load the model

```
In [8]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Conv3D, LSTM, Dense, Dropout, Bidirectional, Ma
        from tensorflow.keras.optimizers import Adam
        from tensorflow.keras.callbacks import ModelCheckpoint, LearningRateScheduler
```

```
In [9]: model = Sequential()
        model.add(Conv3D(128, 3, input_shape=(75,46,140,1), padding='same'))
        model.add(Activation('relu'))
        model.add(MaxPool3D((1,2,2)))

        model.add(Conv3D(256, 3, padding='same'))
        model.add(Activation('relu'))
        model.add(MaxPool3D((1,2,2)))

        model.add(Conv3D(75, 3, padding='same'))
        model.add(Activation('relu'))
        model.add(MaxPool3D((1,2,2)))

        model.add(TimeDistributed(Flatten()))

        model.add(Bidirectional(LSTM(128, kernel_initializer='Orthogonal', return_sequences
        model.add(Dropout(.5))

        model.add(Bidirectional(LSTM(128, kernel_initializer='Orthogonal', return_sequences
        model.add(Dropout(.5))
```

```
model.add(Dense(char_to_num.vocabulary_size()+1, kernel_initializer='he_normal', ac
```

```
In [10]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv3d (Conv3D)	(None, 75, 46, 140, 128)	3584
activation (Activation)	(None, 75, 46, 140, 128)	0
max_pooling3d (MaxPooling3D)	(None, 75, 23, 70, 128)	0
conv3d_1 (Conv3D)	(None, 75, 23, 70, 256)	884992
activation_1 (Activation)	(None, 75, 23, 70, 256)	0
max_pooling3d_1 (MaxPooling3D)	(None, 75, 11, 35, 256)	0
conv3d_2 (Conv3D)	(None, 75, 11, 35, 75)	518475
activation_2 (Activation)	(None, 75, 11, 35, 75)	0
max_pooling3d_2 (MaxPooling3D)	(None, 75, 5, 17, 75)	0
time_distributed (TimeDistributed)	(None, 75, 6375)	0
bidirectional (Bidirectional)	(None, 75, 256)	6660096
dropout (Dropout)	(None, 75, 256)	0
bidirectional_1 (Bidirectional)	(None, 75, 256)	394240
dropout_1 (Dropout)	(None, 75, 256)	0
dense (Dense)	(None, 75, 41)	10537
=====		
Total params: 8,471,924		
Trainable params: 8,471,924		
Non-trainable params: 0		

```
In [11]: model.load_weights('models/checkpoint')
```

```
Out[11]: <tensorflow.python.checkpoint.checkpoint.CheckpointLoadStatus at 0x13e87472050>
```

Build evaluation functions

```
In [12]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
In [13]: def evaluate_performance(actual_texts, predicted_texts):
    # Flatten the lists of strings to lists of characters for character-level metrics
    actual_chars = [char for text in actual_texts for char in list(text)]
    predicted_chars = [char for text in predicted_texts for char in list(text)]

    # Calculate accuracy
    accuracy = accuracy_score(actual_chars, predicted_chars)

    # Calculate precision, recall, and F1 score
    precision = precision_score(actual_chars, predicted_chars, average='weighted')
    recall = recall_score(actual_chars, predicted_chars, average='weighted')
    f1 = f1_score(actual_chars, predicted_chars, average='weighted')

    return accuracy, precision, recall, f1
```

```
In [14]: def add_padding(actual_texts, predicted_texts):
    for i in range(len(actual_texts)):
        len_diff = len(actual_texts[i]) - len(predicted_texts[i])
        if len_diff < 0:
            actual_texts[i] += " " * abs(len_diff)
        elif len_diff > 0:
            predicted_texts[i] += " " * abs(len_diff)
```

Run evaluation

```
In [15]: test_data = test.as_numpy_iterator()
```

```
In [16]: actual_texts = list()
    predicted_texts = list()
```

```
In [17]: for i in range(len(test)):
    print("iteration:", i+1)

    # move to next set
    sample = test_data.next()

    # get real text
    yhat = model.predict(sample[0])
    text_real = [tf.strings.reduce_join([num_to_char(word) for word in sentence]) for sentence in sample[1]]
    text_real = [text_real[0].numpy().decode(), text_real[1].numpy().decode()]
    actual_texts.extend(text_real)

    # get predicted text
    decoded = tf.keras.backend.ctc_decode(yhat, input_length=[75,75], greedy=True)[0]
    text_pred = [tf.strings.reduce_join([num_to_char(word) for word in sentence]) for sentence in decoded]
    text_pred = [text_pred[0].numpy().decode(), text_pred[1].numpy().decode()]
    predicted_texts.extend(text_pred)
```

```
iteration: 1
1/1 [=====] - 4s 4s/step
iteration: 2
1/1 [=====] - 2s 2s/step
iteration: 3
1/1 [=====] - 2s 2s/step
iteration: 4
1/1 [=====] - 2s 2s/step
iteration: 5
1/1 [=====] - 2s 2s/step
iteration: 6
1/1 [=====] - 2s 2s/step
iteration: 7
1/1 [=====] - 2s 2s/step
iteration: 8
1/1 [=====] - 2s 2s/step
iteration: 9
1/1 [=====] - 2s 2s/step
iteration: 10
1/1 [=====] - 2s 2s/step
iteration: 11
1/1 [=====] - 2s 2s/step
iteration: 12
1/1 [=====] - 2s 2s/step
iteration: 13
1/1 [=====] - 2s 2s/step
iteration: 14
1/1 [=====] - 2s 2s/step
iteration: 15
1/1 [=====] - 2s 2s/step
iteration: 16
1/1 [=====] - 2s 2s/step
iteration: 17
1/1 [=====] - 2s 2s/step
iteration: 18
1/1 [=====] - 2s 2s/step
iteration: 19
1/1 [=====] - 2s 2s/step
iteration: 20
1/1 [=====] - 2s 2s/step
iteration: 21
1/1 [=====] - 2s 2s/step
iteration: 22
1/1 [=====] - 2s 2s/step
iteration: 23
1/1 [=====] - 2s 2s/step
iteration: 24
1/1 [=====] - 2s 2s/step
iteration: 25
1/1 [=====] - 2s 2s/step
iteration: 26
1/1 [=====] - 2s 2s/step
iteration: 27
1/1 [=====] - 2s 2s/step
iteration: 28
1/1 [=====] - 2s 2s/step
```

```
iteration: 29
1/1 [=====] - 2s 2s/step
iteration: 30
1/1 [=====] - 2s 2s/step
iteration: 31
1/1 [=====] - 2s 2s/step
iteration: 32
1/1 [=====] - 2s 2s/step
iteration: 33
1/1 [=====] - 2s 2s/step
iteration: 34
1/1 [=====] - 2s 2s/step
iteration: 35
1/1 [=====] - 2s 2s/step
iteration: 36
1/1 [=====] - 2s 2s/step
iteration: 37
1/1 [=====] - 2s 2s/step
iteration: 38
1/1 [=====] - 2s 2s/step
iteration: 39
1/1 [=====] - 2s 2s/step
iteration: 40
1/1 [=====] - 2s 2s/step
iteration: 41
1/1 [=====] - 2s 2s/step
iteration: 42
1/1 [=====] - 2s 2s/step
iteration: 43
1/1 [=====] - 2s 2s/step
iteration: 44
1/1 [=====] - 2s 2s/step
iteration: 45
1/1 [=====] - 2s 2s/step
iteration: 46
1/1 [=====] - 2s 2s/step
iteration: 47
1/1 [=====] - 2s 2s/step
iteration: 48
1/1 [=====] - 2s 2s/step
iteration: 49
1/1 [=====] - 2s 2s/step
iteration: 50
1/1 [=====] - 2s 2s/step
```

```
In [18]: # add whitespace padding for unequal real-predicted pairs
add_padding(actual_texts, predicted_texts)
```

```
In [19]: accuracy, precision, recall, f1 = evaluate_performance(actual_texts, predicted_text

print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1 Score: {f1}")
```

Accuracy: 0.9791415964701163
Precision: 0.9792040812714483
Recall: 0.9791415964701163
F1 Score: 0.9791374960470012